

Assignment 1: Median Redux

Performance comparison of different implementations of MapReduce programs

Pramod Khare

1st February 2015

1 Introduction

There are four implementations of Median Redux:

- (v1) Java sequential without threading/concurrent programming,
- (v2) Map Reduce (sort in Reduce),
- (v3) Map Reduce (using composite keys to let MR do the sort),
- (v4) Map Reduce (either form) where each call to map also computes Fibonacci of N.

2 Note

Performance comparison of multiple MedianRedux implementations is done on following machine hardware Configurations, (Virtual Machine) OS Ubuntu 14.04 LTS, RAM-2.8GB, Processor Core 2 each of Intel Core i3 2.20GHz, No Graphics Processor.

3 Answer Q1

Here are total run-time results from best runs for each version in standalone and pseudo distributed mode:

(time in sec)	V1	V2	V3	V4 (for N <= 10)
Standalone	131.764	272.473	398.346	376.483
Pseudo-distributed	-	319.232	387.790	338.019

From multiple runs of programs v1, v2, v3, it seems clear that all three implementations take approximately similar amount of time. Machine configurations badly affected the total-run-time because of total available physical memory and CPU.

4 Answer Q2

From multiple runs of V4 program, I found out that for larger values of N i.e. more than 10 the run-time of v4 increases significantly.

If we compare v2—v3 with v4, it's pretty clear that for smaller values of N (≤ 10), the time taken by mapper function to compute Fibonacci series is small but negligible and total-run-times are almost the similar.

5 Answer Q3

I found out that the number of reducers is same for all (v2, v3, v4) versions on pseudo-distributed mode i.e. 13.

But for standalone mode, its 51 for all v2, v3, v4 versions.

6 Conclusion

From above comparisons its evident that machine configurations does play a major role in affecting MapReduce performance in pseudo and standalone mode.